



Interactive Body-Driven Graphics for Augmented Video Performance

Nazmus Saquib, Rubaiat Habib Kazi, Li-Yi Wei, Wilmot Li

► To cite this version:

Nazmus Saquib, Rubaiat Habib Kazi, Li-Yi Wei, Wilmot Li. Interactive Body-Driven Graphics for Augmented Video Performance. CHI Conference on Human Factors in Computing Systems Proceedings (CHI 2019, May 2019, Glasgow, United Kingdom. 10.1145/3290605 . hal-02005318

HAL Id: hal-02005318

<https://hal.archives-ouvertes.fr/hal-02005318>

Submitted on 4 Feb 2019

HAL is a multi-disciplinary open access archive for the deposit and dissemination of scientific research documents, whether they are published or not. The documents may come from teaching and research institutions in France or abroad, or from public or private research centers.

L'archive ouverte pluridisciplinaire **HAL**, est destinée au dépôt et à la diffusion de documents scientifiques de niveau recherche, publiés ou non, émanant des établissements d'enseignement et de recherche français ou étrangers, des laboratoires publics ou privés.

Interactive Body-Driven Graphics for Augmented Video Performance

Nazmus Saquib
MIT Media Lab

Rubaiat Habib Kazi
Li-Yi Wei
Adobe Research

Wilmot Li
Adobe Research

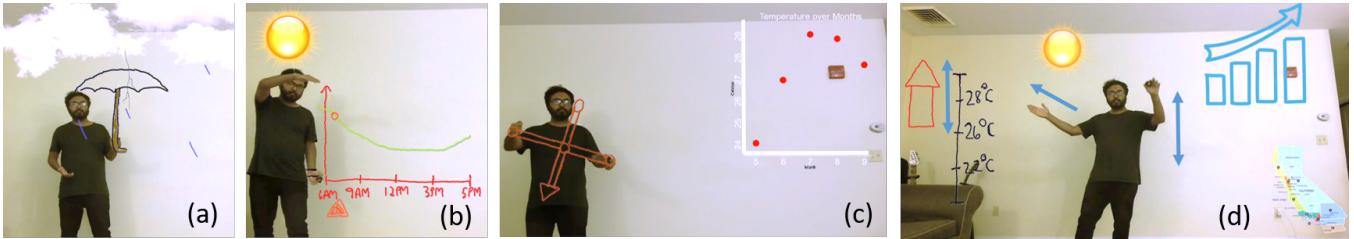


Figure 1: Our system is a presentation tool that lets the performer interactively control the graphical elements using a wide range of hand gestures and body postures for a real-time storytelling experience. (a) An umbrella is attached to the presenter's hand, (b) controlling a custom-made slider to change the humidity chart, (c) directly manipulating a virtual sketch (cross arrow) to change the timeline of a data chart, and (d) pointing and flicking gestures (blue arrows) to move the temperature arrow, which manipulates the sun's opacity.

ABSTRACT

We present a system that augments live presentation videos with interactive graphics to create a powerful and expressive storytelling environment. Using our system, the presenter interacts with the graphical elements in real-time with gestures and postures, thus leveraging our innate, everyday skills to enhance our communication capabilities with the audience. However, crafting such an interactive and expressive performance typically requires programming, or highly-specialized tools tailored for experts. Our core contribution is a flexible, direct manipulation UI which enables amateurs and experts to craft such presentations beforehand by mapping a variety of body movements to a wide range of graphical manipulations. By simplifying the mapping between gestures, postures, and their corresponding output effects, our UI enables users to craft customized, rich interactions with the graphical elements. Our user study demonstrates the potential usage and unique affordance of this mixed-reality medium for storytelling and presentation across a range of application domains.

CCS CONCEPTS

- Human-centered computing → Mixed / augmented reality; Gestural input; User interface design.

ACM Reference Format:

Nazmus Saquib, Rubaiat Habib Kazi, Li-Yi Wei, and Wilmot Li. 2019. Interactive Body-Driven Graphics for Augmented Video Performance. In *CHI Conference on Human Factors in Computing Systems Proceedings (CHI 2019)*, May 4–9, 2019, Glasgow, Scotland UK. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3290605.3300852>

1 INTRODUCTION

"In the coming age, Computer Graphics will become an integral part of our language" - Ken Perlin [41]

Augmented and mixed reality technologies enable user experiences that leverage virtual elements to alter, enhance, and extend our perception of the real world. One simple but powerful form of augmentation is blending animated graphical elements like illustrations, icons, and text with live action footage of performers (Figure 1). Historically, this technique has been used as a special effect for various types of content, including scientific documentaries, instructional material, and music videos. While augmented graphics are typically added as a post-process after the primary footage has been captured, some examples involve performers manipulating graphics in real-time, as in weather forecasts or more recently, social media apps with video overlays.

The expressiveness of augmented videos stems primarily from the range of possible interactions between human performers and animated graphics. To better understand this interaction space, we have analyzed a diverse set of examples and found that augmented graphics are triggered and controlled by a broad spectrum of different gestures and body poses. For instance, graphics are sometimes “attached” to parts of an actor or the scene to highlight, decorate, and enhance the video footage. One specific example of this technique is adding digital clothing or accessories to a character. In other cases, actors directly drive the appearance or motion of graphics by performing a specific pose or gesture, like sweeping their hand to change the scale of an overlaid data visualization or manipulating a virtual, animated object.

While the richness of the interaction space makes it possible to produce a wide range of augmented effects, many of these effects are hard to achieve in practice. Producing real-time augmented videos, like weather forecasts, usually involves specialized professional tools that require a production team to prepare and trigger graphics manually in response to the performer [21, 40]. Video filters on social media apps offer an easier way to achieve real-time effects, but they support a very constrained set of gestures and graphics that limits the expressiveness and potential applications. An alternative is to add augmented effects via video post-processing, using tools like After Effects. However, this approach is restricted to users with the time and expertise to execute advanced video editing and compositing work. Moreover, post-processing is obviously not suitable for real-time augmented videos.

In light of these challenges, we present an authoring tool that helps novice user produce real-time augmented video presentations with a wide range of gestures, body poses, and graphical effects. Our system provides an offline setup mode where users first map gestural actions to shape and appearance attributes of graphical objects via a direct manipulation interface. Then, in the interactive performance mode, the system interprets user actions in real-time, generates the corresponding animated graphics based on the authored mappings, and overlays the graphics onto the video footage. While our method is designed primarily to support real-time scenarios, the proposed approach also facilitates the creation of traditional videos; in this setting, the live performance mode can be used to generate an initial version that users can refine with additional post-processing if necessary.

Our main high-level contribution is in the design of the direct manipulation mapping interface, which provides a flexible, customizable way to associate hand gestures and body postures with graphical effects. We leverage existing hand gesture taxonomies in HCI to represent common gesture categories [13] as well as static body postures as design elements. Using a relational graph structure [30], the user

then maps the variety of gestures and postures to graphical actions – triggering, direct manipulation, indirect parameter tuning, and deformation of the graphical elements. By providing a useful set of predefined and composable primitives, our approach enables users to author a wide range of augmented effects without requiring any explicit programming.

We demonstrate the expressiveness of our system by creating several different styles of augmented video examples, including animated stories, scientific lectures, and explanatory videos. In addition, we have conducted design sessions with a diverse set of users to gain insights about the capabilities and limitations of our tool. We have received positive (albeit preliminary) feedback on the usability, benefits, and potential applications of our approach. The main contributions of this paper include:

- A direct manipulation interface for authoring how input bodies map to output graphical effects.
- An interactive performance interface that applies these mappings in real-time.
- A categorization of gestures and postures based on their different capability and suitability for various mapping scenarios.

2 RELATED WORK

Gestures and HCI

Myron Krueger’s Videoplace [35] is one of the earliest explorations demonstrating a virtual environment that responds to dynamic human gestures for interactions (e.g., painting, pointing, selection). In order to overcome the limitations of WIMP (windows, icons, menus, pointers) interactions, HCI researchers have explored a variety of novel interfaces to leverage the qualities of mid-air gestures for sign-language [44], retrieving and manipulating imaginary objects in 3D [45, 50], and interacting with multi-touch screens or AR (augmented reality) tabletops [16, 32, 46, 48, 49].

In this paper, we consider how to leverage the communicative aspects of gestures to enhance real-time human-to-human communication through dynamic graphics. In this vein, prior systems in HCI and graphics, such as Charade [15], ChalkTalk AR/VR [42], and live multimedia presentation [39], use dynamic hand gestures to control a computer-aided presentation to communicate with the audience. Inspired by this existing work, we propose a direct manipulation interface that enables user to define their own mappings from input actions (gestures and postures) to output graphical effects without programming. Given the idiosyncratic nature of gestures, this flexibility facilitates diverse usage scenarios across many domains. Moreover, our representation and flexible UI accommodates a wide range of input actions (gestures, postures) and output graphical effects, thus enabling an expressive range of interactions with graphical elements.

Interfaces for Dynamic Media and Performance

With the advent of digital technologies, interactive and animated graphics are becoming more popular, prevalent, and a powerful medium for visual art, design, and communication. In general, crafting expressive performance-driven graphical effects requires programming (e.g., openFrameworks [8], processing [9], Unity [10], Flash [7], d3 [6]), highly specialized pre-processing (or rigging) workflows, [12], or timeline-based post-processing tools [4, 5] that require extensive expertise. They are also tailored for application domains (e.g., cartoon characters [21], weather forecasting [40]), with limited interaction capabilities.

In order to make dynamic and animated media accessible, HCI researchers have explored sketch-based and direct manipulation interfaces for animation [23, 27, 28, 31], interface prototyping [36, 38], explanatory illustrations [30, 52], data-storytelling [37], and designing live-performance triggers [47]. Performance based systems have also been explored to map captured human motion into digital characters [24] and arbitrary digital objects (with different topology) [20]. Visual programming tools, such as Scratch [43] and [51], empower users to create their own animations, music, and interactive stories with programmable constructs. Rather than relying on pre-defined models or programming, tools like Kitty [30] and SketchStudio [33] provide an interface where relationships and events can be defined by directly manipulating the underlying relational graph, displayed in the context of the illustration. However, the resulting dynamic artifacts from these tools are designed for pen and touch interactions with limited degrees of freedom. In contrast, our system extends these ideas to produce real-time augmented videos of human performances that leverage the many degrees of freedom of whole body interactions.

Kinect-based Interfaces for Embodied Interaction

With the recent advent of computer vision and real-time pose estimation technologies (Microsoft Kinect, OpenPose [19]), there has been increased interest in exploring how embodied interaction [25] facilitates increased engagement, new learning experiences, and social interactions [29]. In AR mirrors, utilizing the entire body through movement or gesture can support new forms of computer-mediated learning [17, 29] and motor skill improvements [14]. Our work aims to realize these benefits in the context of augmented performance videos for storytelling.

Design Space of Presentation Tools

Since many of the existing systems and techniques discussed above enable users to present ideas with graphics, it is worth examining how our work fits into the overall design space of graphical presentation tools. One way to characterize this

space is to consider two related dimensions: 1) what authors must do to *prepare* the presentation content and 2) what input methods and interactive capabilities are available to *perform* that content. For example, post-processing tools like Adobe After Effects [4] require authors to prepare all of the content by capturing and editing live action footage of performers and then compositing graphical effects. Since the resulting video is created via an offline process, actors cannot interact with the graphics at all during their performance. Our system is more similar to ChalkTalk [41] and performance-driven animation tools that require some amount of preparation to create graphical assets and map them to allowable interactive behaviors and as a result, support real-time interaction with graphics at performance-time. Table 1 summarizes the design space for various presentation methods based on how users prepare and perform content.

	Preparation	Performance input
Slides	Create and layout slides	Trigger events with mouse/clicker
After Effects	Capture and edit footage, composite graphics	-
Chalktalk	Create interactive graphics with programming	Trigger graphics and define mappings with drawn strokes
Character Animator	Create digital characters, map to voice and facial expressions	Trigger and modify graphics with voice and facial acting
Our system	Create graphics, map to gestures and poses	Trigger and modify graphics with gestures and poses

Table 1: Design space for different performance methods.

3 INTERACTION SPACE: AN INFORMAL ANALYSIS

To better understand the types of whole-body interactions humans perform with virtual graphical elements, we analyzed a set of 19 existing augmented videos. We noted the types of input actions, and the resulting graphical effects. Our set of videos included both live and post-processed examples across a range of application domains (e.g., science documentary [1, 3], commercials [2], music videos, and weather forecasting [11]), where the performer interacts with virtual graphical elements to tell stories.

Observed Gestures, Postures, and Parameterization

Within our analysis, we observed a range of (static and dynamic) input actions with high-degrees of freedom, resulting in a wide and rich set of interactions with the virtual elements. Below we discuss our observations, and present a few representative examples.

Hand Gestures. Consistent with prior literature on gesture taxonomies in HCI [13], within our analysis, we identified

a range of gestures types to interact with the *graphical elements*. While such categories of gestures are previously defined in HCI literature [13], in this section we report the interaction capabilities afforded by those gestures, and a range of corresponding output *graphical effects* for storytelling. We refer the readers to Figure 4 in [13] for a nice visual categorization of various gesture types.

Pantomimic gestures are used to mimic an interaction with an imaginary, virtual object [13]. We observed several types of *pantomimic* interaction. First, users often use both hands to manipulate the transformation parameters (translation, rotation, or scale) of the virtual object (e.g., rotating a rigid mug, photo frame). Second, pantomime gestures can also be used to deform the shape of a graphical object (e.g., interacting with a joystick, or an octopus). In such cases, the overall transformation parameters remain the same.

Iconic gestures are used to communicate information about objects or entities, such as specific sizes, shapes, and motion paths. We also noted *pointing* to highlight, and a variety of *direct manipulation* techniques.

Semaphoric gestures are hand movements and postures that convey specific meanings. *Semaphoric strokes* represent hand flicks which are single, stroke-like movements [13]. We observed that *semaphoric strokes* are used to manipulate graphical parameters, which are challenging to communicate using *direct manipulation* or *pantomimic* gestures, due to proximity, scale, or abstraction. For instance, a science presenter repeatedly flicks one hand to communicate that air is vacuumed out of a jar while a vacuuming animation is displayed. In this case, the *semaphoric stroke* gesture is used to manipulate a parameter in a more abstract way. Such gestures are also used to point a graphic at a distant location (e.g., a timeline) and manipulate, to counteract the scale and proximity of the graphical element. Since *semaphoric* gestures are strictly learned, such mappings between the gesture and the parameters should be pre-defined by the user.

Rigging and Static Body Postures. Static and dynamic *body postures* are also used to interact with the graphical elements. Similar to gestures [15, 18, 45], there are a number of advantages of using body postures to interact with virtual *graphical elements*. A static or dynamic posture can be used to specify both a command and its parameters. For instance, a guitar holding posture would indicate the users intention to trigger the guitar graphics. Further, the posture information (e.g., the position of the hands) also specifies the parameter of the guitar graphics, including the position, orientation, and scale. Often, virtual sketches and graphics (clothes, helmet, other wearables) are rigged and anchored with respect to the human(s) in the video, as well as arbitrary objects in the

scene. In such cases, the sketches are deformed dynamically in accordance to the human skeleton.

Insights and Observations

This body-driven medium affords a rich interaction space, in terms of input actions (body postures, skeleton, and variety of gestures) and graphical output effects (direct and indirect manipulations, deformation, spatial, visual, temporal, and quantitative parameter tuning).

The types of gestures observed are consistent with gesture taxonomies in HCI. We also found postures and skeletal rigging as a powerful elements to interact with graphics.

Depending on the context and applications, a gesture could be interpreted in multiple ways. For example, the gestures and postures used by a musical band can be very different from an educational setting. Hence, instead of predefining the effects of gestures, we should enable users to map their input actions into output effects.

Despite the potential benefits of this medium, there is a lack of tools that leverage the richness of the interaction space and facilitate flexible mapping between input actions to output effects to a wider range of audience.

Design Goals. Based on our observation and analysis, we formulate the following design goals for our system:

- (1) A direct-manipulation UI that is intuitive to use, and yet comprehensive to achieve diverse mapping of human actions to graphical effects without the complications of programming or traditional video editing.
- (2) A flexible rigging system that is simple, expressive, and facilitates a wide range of parameterizations to interact with the graphical elements.
- (3) Providing interactive controls during live performance to address some of the challenges of a gesture based system, such as ambiguity, segmentation, tracking accuracy, and visual feedback.

4 USER INTERFACE AND INTERACTIONS

We have designed and implemented a system that enables users to craft augmented, interactive live-streaming presentations. Our system has an offline authoring environment to craft a story, where the presenter prepares the (static and animated) graphical elements, and maps gestures and postures to the parameters of graphical elements. Figure 2 shows our user interface, consisting of a main canvas, a reference skeleton representing the presenter, as well as global and contextual toolbars. In addition to offline authoring, our system also supports a live performance environment that allows the presenter to perform and interact with the graphical elements in real-time.

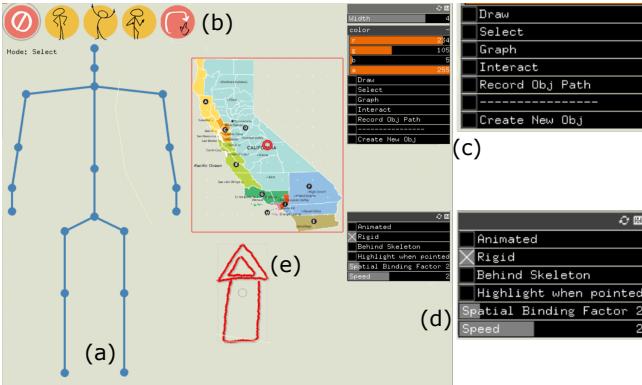


Figure 2: Our system authoring UI. (a) Reference skeleton, (b) Trigger menu consisting of template postures and iconic gesture, (c) Global toolbar for drawing and system modes, (d) Contextual toolbar for graphical elements, (e) Graphical elements: an imported map and a sketch.

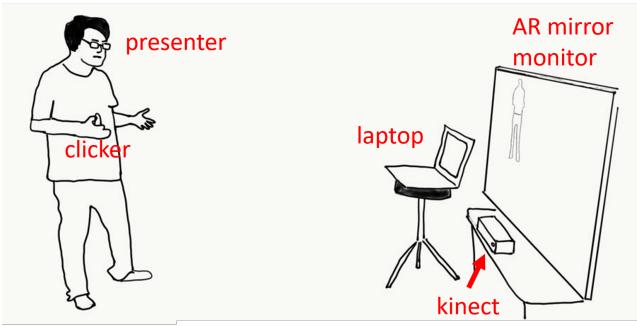


Figure 3: Hardware setup.

Hardware setup

Our current setup consists of a skeleton tracking device (Kinect) facing the user, a computer for authoring and playback, and a wearable clicker for interactive controls during live presentation (Figure 3). An optional large monitor serves as an AR (augmented reality) mirror to facilitate better viewing and feedback for the presenters.

Creating and Importing Graphical Elements

Our system supports a variety of graphical elements to facilitate expressive storytelling: sketches, images, animated GIFs and textures, and 2D scatter plots.

To create a *sketch element*, the user sketches a few strokes using the *pen tool*, and clicks the *create element tool*. This aggregates all the strokes into a single *graphical element*. The *element* contextual toolbar allows the user to further configure the graphical element. The user can sketch a *translation path* [30] to define a parameterized polygon path, which constrains the movement of the *graphical element* along the

path during live presentation mode. The user can also configure the *anchor point* of the *element* by direct manipulation. Clicking the *animated* button in the contextual toolbar turns the *sketch element* into an *animated texture* [30]. The user can also directly drag and drop external static images, animated GIFs, and data files (in CSV format for customizable 2D charts) to create graphical elements. Using the *selection tool*, the user can select, move, edit, and delete a *graphical element* any time. Some of these effects are shown in Figure 1.

Input Actions: Skeleton, Gestures, and Postures

During the authoring mode, a reference skeleton (Figure 2) represents the performer. In the design of our system, a key challenge is to categorize and represent the high-dimensional input actions into meaningful, concise entities. To this end, we leverage Aigner et al.'s existing gestural taxonomy in HCI [13] to represent the gestural actions - *pantomimic*, *direct manipulation*, *semaphoric*, *pointing*, and *iconic*. We also provide a customizable template of *static postures* to interact with the graphics (Figure 2). These postures can trigger the *graphical elements*, and set their positions, scales, and orientations. The *skeletal joints* can also be used to drive the position and deformation of *graphical elements* by anchoring and rigging. Overall, the variety of gestures, postures, and skeletal joints accommodate a range of input actions to interact with the *graphical elements*.

Output Effects: Parameterization and Deformation

Figure 4 shows the types of parameter manipulations a *graphical element* can undergo in our system, including visual, spatial, temporal, quantitative parameters, as well as causal relations and freeform deformation. We normalize the range of each parameter value from 0 to 1. As for non-rigid *graphical elements*, the user can define multiple *pins* (or constrained handles) within the bounding box. During live presentation, the position of one or more of those *pins* are driven by the *skeletal joints* of the presenter for non-rigid deformations. As demonstrated in our results, the range of parameters and flexible riggings facilitates rich interaction possibilities.

Mapping: Input Actions to Output Effects

In this section, we will discuss the interactions that map input actions (*gestures*, *postures*, *skeletal joints*) to output effects (*direct* and *indirect manipulations*, *deformation*, and *anchoring*). Such mappings are achieved by creating *edges*, or by associating a gesture or posture directly to a *graphical element*. Before we dive into individual interaction techniques, we describe the underlying relational graph structure that captures such mappings and causal relationships in the scene.

Relational Graph. Central to all interactions in our system is a graph data structure [30] that treats all *graphical elements*

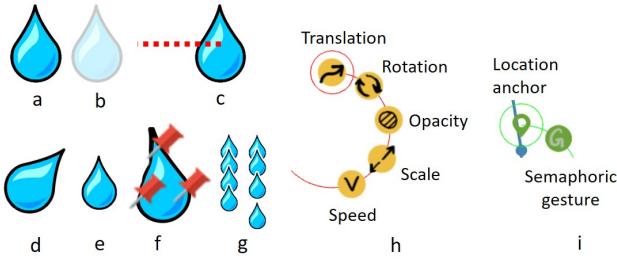


Figure 4: Possible parameters for interacting with a graphical object. (a) original graphic, (b) changing opacity (visual parameter), (c) translating along a user defined path, (d) rotation, (e) scaling, (f) non-rigid deformation using control points (pins), (g) speed of animation (temporal), (h) graphical element radial menu, displaying the parameters, (i) skeleton joint radial menu.

(sketches, images, GIFs, animated texture) and the *skeleton* as *nodes*. The *edges* in the graph define the mapping and functional relationship between the source and destination node parameters, thus depicting the coordination and causal relations between them. In addition, *edges* also capture the rigging mechanism of our tool for deformation, from *skeletal joints* to *graphical element pins*. Users can view the graph structure (nodes and edges) by clicking the *Graph* button in the main toolbar. Figure 5 shows an example.

Creating Edges for Causal Effects. To create a causal relationship between two *graphical elements*, the user sketches an edge from source to destination node. Similar to Kitty [31], a contextual *radial menu* and *functional relationship* widget enable users to specify the source node attribute (Figure 6), destination node attribute, and the corresponding parameterization function (via sketching), where the horizontal and vertical axes represent the driver and driven parameters.

Edges for Direct Manipulation and Pantomimic gestures. Our system supports uni-manual and bi-manual direct manipulation interaction with the virtual *graphical elements*. We discuss a few *direct manipulation* techniques using our system. To anchor a *graphical element* to a *skeletal joint* (hand, for instance), the user switches to *Graph* mode, and then sketches an *edge* from the hand *joint* to the *graphical element pivot*. This edge creates a new mapping. During live performance, when the element appears in the scene, the *graphical element* follows the *skeletal joint* freely (Figure 5d→e). However, the movement of a *graphical element* can also be constrained by a parameterized, user-defined *translation path* (Figure 6). Such constrained manipulations are useful to create custom sliders (e.g., timeline) and other forms of interactions.

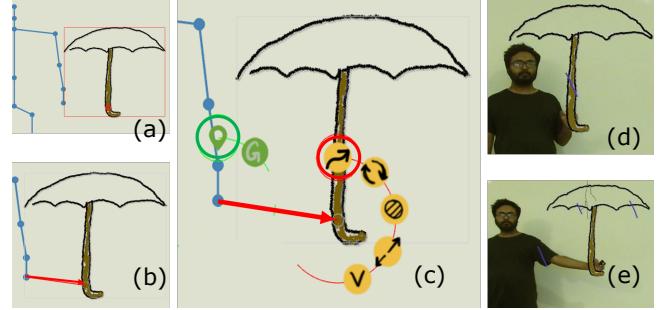


Figure 5: Creating an edge in the relational graph structure to anchor a graphical element (umbrella) to the hand joint. (a, b) Dragging a line in graph mode to create an edge between a hand joint and the umbrella. (c) Radial menus for a skeleton joint and a graphical element in the graph mode. (d, e) In live playback, the umbrella is anchored at the hand by its pivot point and moves freely.

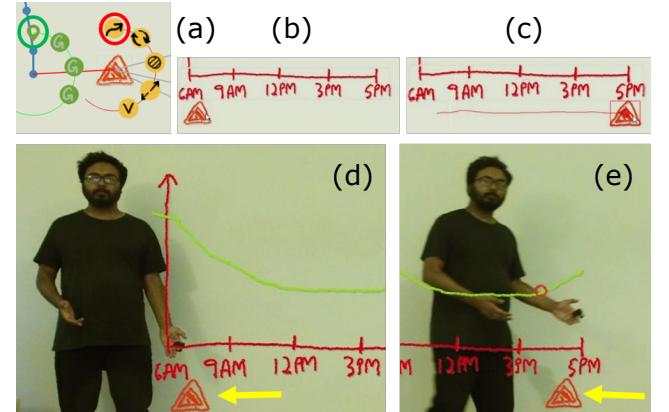


Figure 6: Defining a path to make a custom slider. (a) Creating an edge between a hand and the graphical element for slider. (b, c) Dragging the slider graphic in the desired direction in path mode. (d, e) The graphical element moving along a constrained path following the hand (yellow arrow).

As for *pantomimic gestures*, the user interacts with a single *graphical element* by grabbing it with both hands. In such cases, during the authoring mode, the user defines multiple pins in the *graphical element*, and then specifies one edge from each hand joint to the corresponding *pin*. During live performance, our system computes the optimum rigid transformation parameters (position and rotation) of the element that best matches the hand positions (Figure 7). This setting allows the user to manipulate a rigid *graphical element* in a way to how *pantomimic gestures* are used. The transformation parameters of the *graphical element* can also drive other elements in the scene using causal-and-effect relationships (e.g., by connecting the rotation parameter of the element to a parameter of another element).

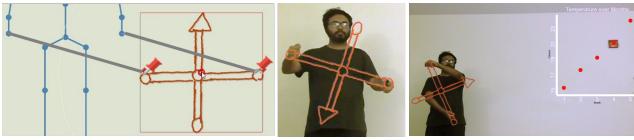


Figure 7: Pantomimic gesture example. Left: Pins on graphical element and corresponding edges from skeleton joints. Middle: Pantomimic gesture for rigid transformation. Right: The rotation parameter drives the data-chart timeline.

Iconic gestures. An *iconic gesture* specifies the size and position of a *graphical element* when it appears in the scene (Figure 8). For such mappings, during the authoring mode, when a *graphical element* is selected, our system displays a *trigger menu* at the top of the canvas (Figure 2(b)). This menu specifies variety of behaviors that take effect during the triggering of the *element*. The user then selects *iconic gesture* as the triggering behavior. Unlike other gesture types, iconic gestures take effect only when the element appears in the scene for the first time. This is why we display this gesture in the *trigger menu*, as opposed edges for manipulation. During live presentation, the user performs a freeform gesture, whose position and size sets the position and size of the *graphical element*. As Aigner et al. reported [13], *iconic gestures* are spontaneous, and do not rely on commonly known vocabulary. Thus, we do not take the actual shape into account. Only the position and size of the gesture matters for the parameterization.

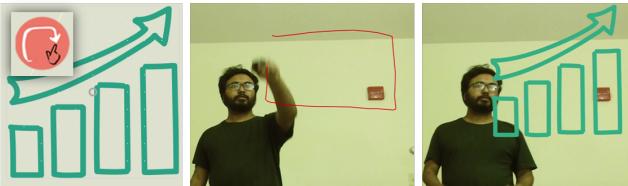


Figure 8: Iconic gesture example. Left: defining the trigger by selecting the gesture (inset). Middle: performing the gesture in live playback (our system draws a trace of the gesture performance). Right: graphic element is positioned and scaled according to where and how big was the gesture.

Pointing. *Pointing* is used to highlight and select a *graphical element*. The user can activate the "pointing" option using a check box in the *element's* contextual menu, so that during the live performance, when the performer points to that element, it is selected and highlighted. By default, this option is de-activated for the *elements*.

Edges for semaphoric gestures. Unlike *direct manipulation* techniques described above, *semaphoric gestures* in our system are used to make indirect manipulation of graphical

parameter using flick gestures (Figure 9). To map a parameter of a graphical element to a *semaphoric gesture*, the user switches to *graph mode*, and creates an *edge* from the hand joint to the *element*. The user then selects the *semaphore* attribute from the *element* radial menu. This creates a mapping between the *semaphoric gesture* to the corresponding attribute. Such gesture interactions are suitable to manipulate abstract parameters, as well as interacting with distant elements.



Figure 9: Semaphoric gesture example. When the user waves the hand, the temperature arrow moves up. Blue arrows are added to visualize the direction of pointing and movement.

Edges for rigging and deformation. Our system's flexible rigging and mapping also supports a variety of interaction capabilities with deformable graphics that can be controlled by the user's body. In order to create such mapping, the user sets the *graphical element* as non-rigid (via contextual toolbar), and defines a few *pins* that serve as constrained handles [28]. The user then maps the pins to skeletal control points by sketching *edges* between them. Figure 10 demonstrates the process of creating deformable interactions.

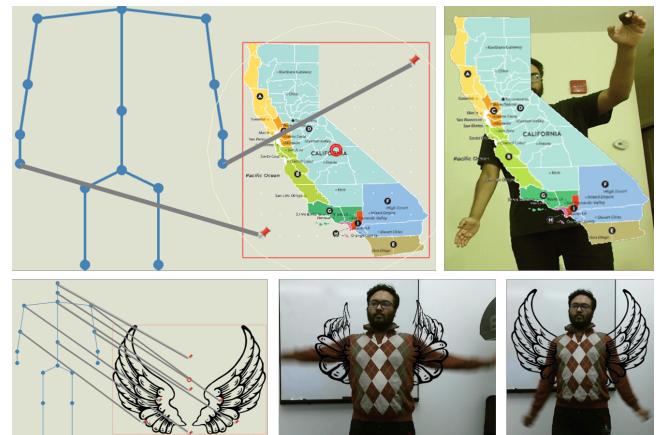


Figure 10: Deformation examples. Top: manipulating a map by two hands. Bottom: flapping wings via arm motions.

Body Posture. Our system allows the user to trigger and parameterize *graphical elements* using pre-defined *body postures*. For instance, when the performer does a guitar holding

posture during the presentation, it triggers the virtual lighting in the scene, and places the element between the two hand positions (Figure 11). To associate a graphical element to a pre-defined body posture, the user selects the corresponding *posture* icon from the trigger menu (Figure 2(b)) associated to the *element*. During the performance, our system detects the pre-defined postures, and trigger the graphics accordingly. Once the performer moves to a different posture, the *graphical element* disappears from the scene. As noted by Steins et al. [45], such postures enable users to quickly switch between devices to optimally support the current task. They also exploit users' knowledge and experience using physical devices, which makes such interaction self-revealing. We provide three postures that were widely observed during our informal study for highlighting or triggering objects. However, depending on the use case and applications, the user can train the system for their own *posture* templates.

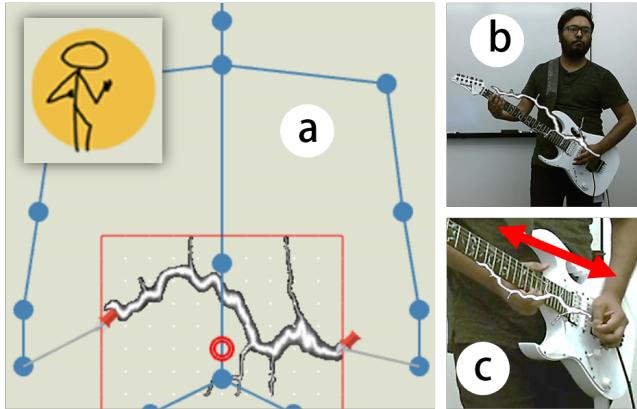


Figure 11: Static pose trigger and subsequent parameterization by non-rigid deformation. (a) Selection of posture trigger (inset, guitar pose), placement of control points on a GIF of a thunder, and rigging with both hands for deformation. (b) The thunder graphic element appears when the guitarist holds the guitar (makes a guitar pose), and (c) it deforms based on where the hands are when the guitarist plays (red arrow added for emphasis).

Story and Scenes

The resulting story consists of a sequence of *scenes* (or slides). A scene consists of a sequence of *graphical elements*, along with the underlying *graph structure*.

Other Features

Spatial binding. Our system enables users to set the *spatial binding* parameter for each *element* via a contextual toolbar. This parameter allows the users to define the distance at which the *elements* will interact with their bodies.

This customization for each *element* is important for a highly interactive, gesture-based system, because the users might not intend for all element to interact at once as they move around or perform a gesture.

Interactions during Live Performance

During live performance, the author presents the interactive story using gestures and postures. However, there are a number of practical challenges in the design of such a body-based system, such as accuracy of gesture recognition systems, the problem of segmentation and chunking [15], and ambiguities. During the performance, the user may want to use gestures and postures freely without impacting the *graphical elements*, even if they are defined as actions and triggers during the authoring mode. The user should have full control over when they want an object to be triggered or interact with existing graphics. This is an ambiguous setting, and defining such nuanced interactions is complex.

To mitigate such problems, we equip the performer with a consumer-grade wearable presenter (known as clicker) to provide controls and interact with the system. We mapped the buttons of the clicker to various system functionalities. The *next* button loads the graphical elements sequentially, while double tapping loads the next *scene*. To perform an *iconic gesture*, or trigger an *element* with a *body posture*, the user can press and hold the lower button. They can also press the upper button to perform a *semaphoric gesture*. Our system only tracks for gestures and postures when the user holds down these buttons. The user can also freeze a scene by disabling *spatial binding* for all *graphical elements* in the scene by double clicking the upper button.

Our system also provides visual cues to aid the performer, such as the next *graphical element* at the top left corner, and a freeze icon (top right) to indicate interaction freeze.

5 METHOD AND IMPLEMENTATION

Our system is developed using Openframeworks, a C++ framework for graphical applications. We implemented our system on an Alienware laptop with Intel i9-8950 processor, 32 GB RAM, and an NVIDIA GTX 1080 Ti graphics card.

Skeleton, Gesture, and Pose Tracking. We have used Kinect for Windows V2 for our hardware setup, producing a frame rate of about 60 FPS with the other processes, such as rendering elements, gesture and posture classifiers. While OpenPose [19] produced a lower frame-rate (8 FPS) on average. We project the Kinect skeleton to the 2D screen, and use an adaptive naive Bayes classifier (ANBC) [26] for static pose recognition (with skeleton joint angles as features).

Graphical Transformations. For non-rigid deformation of images, textures, and sketches, we use the as-rigid-as-possible mesh deformation algorithm with control points [28]. For

	Background	Graphical Tools	Presentation Skills
P_1	2D & 3D artist	Advanced	Intermediate
P_2	Videographer	Expert	Intermediate
P_3	2D Artist	Intermediate	Beginner
P_4	Business, Science	Beginner	Expert
P_5	Content Writer	Intermediate	Expert
P_6	Astrophysicist	Beginner	Advanced

Table 2: Background of our study participants. Four categories of increasing expertise: beginner, intermediate, advanced, and expert.

rigid transformation of objects (*pantomimic gesture*), we calculate the optimum rotation and position of a rigid graphical element based on its control points' average angle and position difference from skeleton joints that are attached to them in our relational graph structure.

6 DESIGN SESSIONS

We conducted informal, qualitative design sessions with a diverse set of users to gain insights about the potential applications, usability, and limitations of our tool. Users were invited to create a presentation of their own using our tool. The sessions lasted between 2 - 3 hours.

Participants

We invited 6 participants (3 females, age range 23 - 35) with diverse backgrounds and professions (Table 2). Such diverse subjects are better suited to evaluate an emerging mixed medium creation tool, which can be used in many settings. Based on their background survey, we have categorized our participants in four categories in increasing expertise: beginner, intermediate, advanced, and expert.

Methodology

Since there is no existing tool that can provide similar live performance authoring in a flexible setting, we did not compare against a baseline case. The study consists of the following three steps.

Training. We introduced the participants with the overall project, and described the functionalities and usages of our system through a guided training session. This gave the participants an opportunity to systematically learn about each component of the system.

Target task. In this step, participants were given a target story to reproduce with our system. A part of the weather report story (Figure 1) was chosen for this, which includes *iconic gestures*, a few graphical transformations based on *direct manipulation, deformation*, a combination of *pointing* and *semaphoric gestures*, and a variety of *graphical elements* (static pictures, textures, and a CSV data file). They were also

asked to perform in the live playback mode to get a sense of how they can interact with the elements present in a scene.

Creating Story and Performance. We then let each participant author their own story. Clarifications were provided if they asked questions about the system. After the authoring and performance, we asked them to fill out a questionnaire.

7 RESULTS

Participant feedbacks. Our participants responded positively to the unique affordances, usability, and novel interaction capabilities of our system.

P_3 : “Each function has a clear goal (like graph, draw) and it is easy to combine them to create different effects.”

However, in contrast to P_{1-3} (artists), the remaining participants (P_{4-6}) were less experienced with graphical tools (animation, video editing). Not surprisingly, some of our system’s concepts and functionalities (rigging, transformation) were completely new to them, requiring more time to familiarize. However, once they were familiarized, their timing for producing the target task was comparable (P_4) or faster ($P_{5,6}$) than the artists. They also felt more comfortable in the live performance mode.

P_6 : “There are some quirks in the system and it takes a bit of time to get used to it. But the learning curve relatively flat and I was able to pick it up rather quickly.”

The participants also responded positively to the unique live storytelling aspects of this medium with interactive graphics. In terms of traditional presentation tools, P_4 commented: “This system gives me much more control over HOW I want to present and WHERE I want to focus the audience’s attention. Storytelling is often about timing and the right level of emphasis, which is much easier to pull off with this than a slide deck.” As for user engagements and applications, P_6 commented - “I engage in STEM outreach quite often and I see potential for software like this to engage with the public. This system allows scientists to transform boring plots to interactive activities and brings the potential to educate the public in a more effective way.”

In contrast to existing video editing tools, participants ($P_{2,5}$) found our system to be easy, flexible, and more accessible, with the added benefit of live-performance and interactivity. P_2 also mentioned the potential of such interfaces and interactions in the context of post-processing video editing tools for greater efficiencies.

Target task. All the participants were able to complete the task without any assistance. The time to reproduce the example in the authoring environment varied from 5 to 14 minutes, with an average of 8.3 minutes. Our program encountered several bugs during P_1 ’s session. Several participants $P_{1,2}$ were confused about the use of *semaphoric gesture* to manipulate timeline using a *graphical element* (at a distant location).

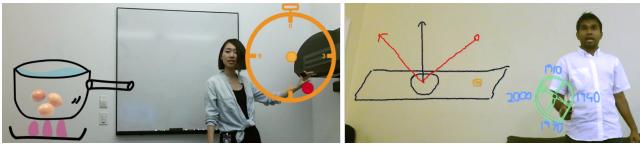


Figure 12: Example stories made by users. (Left) cooking tutorial, (right) presenting an astronomy research paper.

Such confusions could be alleviated by better training and visual feedback during live performances.

Resulting stories in freeform stage. Participants crafted a range of stories and presentations using our system, ranging from a cooking tutorial to a full research paper presentation (Figure 12). Some participants $P_{5,6}$ created stories pertaining to their own professional background, while others were more experimental. Before authoring the presentation, all the participants prepared the storyline, and gathered the necessary *graphical elements* from the web. After the authoring, we recorded the participants live performance presenting the story. Overall, this session took between 1 to 2 hours.

P_1 prepared a story for his youtube channel, using bi-manual deformation, uni-manual *direct manipulation*, *anchoring*, and graphical overlays. P_2 demonstrated how to cook ramen eggs, with 3 scenes, leveraging *direct manipulation*, *causal relationships*, and *iconic gestures*, demonstrating the steps of activities. P_3 performed a presentation about how to meditate in 1 minute.

P_4 's story about interior planning took the longest time to author (1.5 hours). She explored and tested some interactions before choosing a few for her story. Her performance incorporated a number of interactions, including *sempahoric gestures* to clean up the window, *direct manipulation*, and *iconic gesture* to place paintings. She also leveraged her ankle joint to interact with the graphical elements (moving a sofa). P_5 produced a story on a social experiment, it took him 30 minutes to author and test his story before performing.

P_6 had the longest performance among all (13 minutes). He presented an astronomy research paper using our system. The story creation process took an hour overall, and included finding and cleaning a dataset from his computer, and storyboarding to figure out the contents of each slide in our system. He used a *pantomimic gesture* to rotate a time wheel and control a data chart, and *direct manipulation* to slide a night sky image to begin and end his story.

Overall, participants used a variety of gestures and parameterization capabilities, indicating that they found those interactions useful in their performances (Figure 13). Based on the study, we believe the main benefit in supporting a diverse set of gestures is that users are free to define mappings that work well with their specific content or style. They can also choose to limit the range of gestures to reduce cognitive

load. The distribution of gestures in the study largely aligns with this strategy; overall, participants covered all categories of gestures, but within a presentation, they each used a small set. As for critical gestures, the most popular were pointing and manipulation, but all types were used at least 2 times.

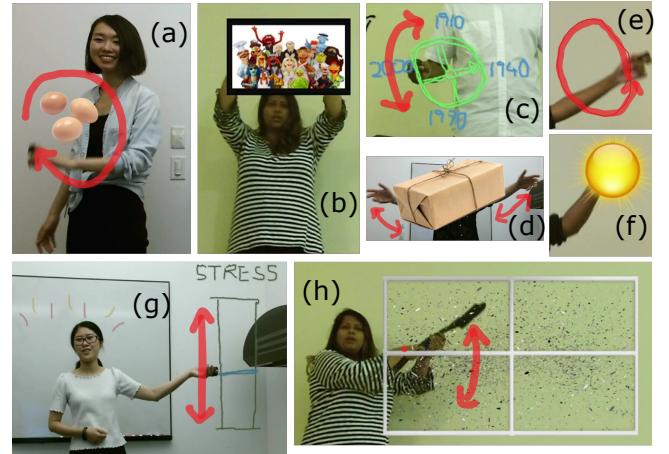


Figure 13: Examples of interactions used by users. (a) P_2 Triggering eggs image using iconic gesture, (b) P_4 triggering an image by a static pose, (c) P_6 pantomimic gesture to create a timeline wheel, (d) P_1 non-rigid deformation of box attached to both hand joints, (e) - (f) P_5 iconic gesture to bring up sun, (g) P_3 direct manipulation (stress meter) and animation anchoring (rays coming out from head), and (h) P_4 semaphoric gesture to clean window (by parameterizing the opacity of dirty and clean window layers).

8 LIMITATIONS AND FUTURE WORK

Despite the encouraging results, there are also a number of limitations and opportunities for future improvements that warrant discussion. Based on the feedbacks from the participants and our own experiences with the system, we plan to investigate several future works:

In our current prototype, the mappings are defined at authoring time that can be applied interactively in an expressive but predictable way at performance time. A potential future work is to allow interactive mapping of graphical effects during performance similar to ChalkTalk [42].

As pointed out by our participants, a major challenge for live performance is the mental overload of the presenter. Better visualization is required to preview or remind about what they need to perform next, and what the graphics effects look like in real-time. In our current implementation, we provide a large display in front of the presenter in the form of an augmented reality mirror (Figure 3). We believe wearable mixed-reality eyeglasses or headsets, equipped with presenter's view, can reduce such mental overloads. In addition to

live presentation, our interface design and system components can also be recombined for video post processing to reduce manual workload.

Our current design and implementation tracks a single skeleton only. We plan to incorporate multiple users to facilitate collaborative live performance with interactive effects.

The current interaction scope is limited by computer vision and machine learning components, such as accuracy of posture and gesture tracking. In the future, we would like to incorporate additional sensing modalities (e.g., Leap Motion, Project Soli) for more fine-grained hand gesture recognition and interactions. Furthermore, by incorporating scene geometry (e.g. via ARCore/ARKit), we can render the graphical elements in planar surfaces (e.g., walls, floors).

Our system can be extended for multi-modal interaction, such as speech in addition to hand gestures and body postures. Inspired by existing videos such as [22, 34], we intend to add physical objects via object recognition, tracking, and segmentation.

9 CONCLUSIONS

Gestures and postures are an integral part of human to human communications. As mixed and augmented reality are becoming widespread, we explore how such natural, innate human capabilities can augment our live storytelling. We designed and implemented a flexible, intuitive UI that enables users to map a range of input actions to output effects, resulting real-time, rich interaction capabilities with graphical elements. The diverse story and performances by our design session participants demonstrate the potential of this interactive, rich, augmented storytelling medium.

REFERENCES

- [1] 2012. The Inexplicable Universe: Unsolved Mysteries. <https://www.netflix.com/title/70305069>. Accessed: 2019-01-07.
- [2] 2012. Microsoft: Your Potential, Our Passion. https://www.youtube.com/watch?v=it-Vt2M_-ZU. Accessed: 2019-01-07.
- [3] 2013. Hans Rosling: The River of Myths. <https://www.youtube.com/watch?v=lYpX4l2UeZg>. Accessed: 2019-01-07.
- [4] 2018. Adobe After Effects. <https://www.adobe.com/products/aftereffects.html>. Accessed: 2018-09-20.
- [5] 2018. Adobe Animate. <https://www.adobe.com/products/animate.html>. Accessed: 2018-09-20.
- [6] 2018. D3.js, Data-Driven Documents. <https://d3js.org/>. Accessed: 2018-09-20.
- [7] 2018. Flash with ActionScript 3.0. <https://www.adobe.com/devnet/actionscript/documentation.html>. Accessed: 2018-09-20.
- [8] 2018. openFrameworks, C++ toolkit for creative coding. <https://openframeworks.cc/>. Accessed: 2018-09-20.
- [9] 2018. Processing. <https://processing.org/>. Accessed: 2018-09-20.
- [10] 2018. Unity. <https://unity3d.com/>. Accessed: 2018-09-20.
- [11] 2018. The Weather Channel's new storm graphics are totally insane. <https://www.fastcompany.com/90176637/the-weather-channels-new-storm-graphics-are-totally-insane>. Accessed: 2019-01-07.
- [12] Adobe. 2018. Adobe Character Animator. <https://www.adobe.com/products/character-animator.html>.
- [13] Roland Aigner, Daniel Wigdor, Hrvoje Benko, Michael Haller, David Lindbauer, Alexandra Ion, Shengdong Zhao, and JTKV Koh. 2012. Understanding mid-air hand gestures: A study of human preferences in usage of gesture types for hci. *Microsoft Research TechReport MSR-TR-2012-111 2* (2012).
- [14] Fraser Anderson, Tovi Grossman, Justin Matejka, and George Fitzmaurice. 2013. YouMove: enhancing movement training with an augmented reality mirror. In *Proceedings of the 26th annual ACM symposium on User interface software and technology*. ACM, 311–320.
- [15] Thomas Baudel and Michel Beaudouin-Lafon. 1993. Charade: remote control of objects using free-hand gestures. *Commun. ACM* 36, 7 (1993), 28–35.
- [16] Hrvoje Benko, Ricardo Jota, and Andrew Wilson. 2012. MirageTable: freehand interaction on a projected augmented reality tabletop. In *Proceedings of the SIGCHI conference on human factors in computing systems*. ACM, 199–208.
- [17] Tobias Blum, Valerie Kleeberger, Christoph Bichlmeier, and Nassir Navab. 2012. mirracle: An augmented reality magic mirror system for anatomy education. In *Virtual Reality Short Papers and Posters (VRW), 2012 IEEE*. IEEE, 115–116.
- [18] William Buxton, Mark Billinghurst, Yves Guiard, Abigail Sellen, and Shumin Zhai. 2002. Human input to computer systems: theories, techniques and technology. *Manuscrito de livro em andamento, sem editora* (2002).
- [19] Zhe Cao, Tomas Simon, Shih-En Wei, and Yaser Sheikh. 2017. Realtime Multi-Person 2D Pose Estimation using Part Affinity Fields. In *CVPR*.
- [20] Jiawen Chen, Shahram Izadi, and Andrew Fitzgibbon. 2012. KinETre: Animating the World with the Human Body. In *Proceedings of the 25th Annual ACM Symposium on User Interface Software and Technology (UIST '12)*. ACM, New York, NY, USA, 435–444. <https://doi.org/10.1145/2380116.2380171>
- [21] CNet. 2016. How Cartoon Donald Trump comes to life on 'The Late Show'. <https://www.cnet.com/news/cartoon-donald-trump-late-show-stephen-colbert/>.
- [22] CNN Money. 2018. This is what a trade war looks like. <https://youtu.be/VA-LdvH35Uk>.
- [23] Richard C Davis, Brien Colwell, and James A Landay. 2008. K-sketch: a'kinetic'sketch pad for novice animators. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. ACM, 413–422.
- [24] Mira Dontcheva, Gary Yingve, and Zoran Popović. 2003. Layered Acting for Character Animation. In *ACM SIGGRAPH 2003 Papers (SIGGRAPH '03)*. ACM, New York, NY, USA, 409–416. <https://doi.org/10.1145/1201775.882285>
- [25] Paul Dourish. 2004. *Where the action is: the foundations of embodied interaction*. MIT press.
- [26] Nicholas Gillian, R Benjamin Knapp, and Sile O'Modhrain. 2011. An adaptive classification algorithm for semiotic musical gestures. Citeseer.
- [27] Robert Held, Ankit Gupta, Brian Curless, and Maneesh Agrawala. 2012. 3D puppetry: a kinect-based interface for 3D animation.. In *UIST*. Citeseer, 423–434.
- [28] Takeo Igarashi, Tomer Moscovich, and John F. Hughes. 2005. As-rigid-as-possible Shape Manipulation. *ACM Trans. Graph.* 24, 3 (July 2005), 1134–1141. <https://doi.org/10.1145/1073204.1073323>
- [29] Seokbin Kang, Leyla Norooz, Vanessa Oguamanam, Angelisa C Plane, Tamara L Clegg, and Jon E Froehlich. 2016. SharedPhys: Live Physiological Sensing, Whole-Body Interaction, and Large-Screen Visualizations to Support Shared Inquiry Experiences. In *Proceedings of the The 15th International Conference on Interaction Design and Children*. ACM, 275–287.

- [30] Rubaiat Habib Kazi, Fanny Chevalier, Tovi Grossman, and George Fitzmaurice. 2014. Kitty: Sketching Dynamic and Interactive Illustrations. In *UIST ’14*. 395–405. <https://doi.org/10.1145/2642918.2647375>
- [31] Rubaiat Habib Kazi, Fanny Chevalier, Tovi Grossman, Shengdong Zhao, and George Fitzmaurice. 2014. Draco: Bringing Life to Illustrations with Kinetic Textures. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI ’14)*. 351–360. <https://doi.org/10.1145/2556288.2556987>
- [32] Rubaiat Habib Kazi, Kien Chuan Chua, Shengdong Zhao, Richard Davis, and Kok-Lim Low. 2011. SandCanvas: a multi-touch art medium inspired by sand animation. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. ACM, 1283–1292.
- [33] Han-Jong Kim, Chang Min Kim, and Tek-Jin Nam. 2018. SketchStudio: Experience Prototyping with 2.5-Dimensional Animated Design Scenarios. In *DIS ’18*. 831–843. <https://doi.org/10.1145/3196709.3196736>
- [34] Yoshikage Kira. 2018. Transportation history. <https://youtu.be/IB3pdb3yXuY>.
- [35] Myron W Krueger, Thomas Gionfriddo, and Katrin Hinrichsen. 1985. VIDEOPLACE—an artificial reality. In *ACM SIGCHI Bulletin*, Vol. 16. ACM, 35–40.
- [36] James A Landay and Brad A Myers. 1995. Interactive sketching for the early stages of user interface design. In *Proceedings of the SIGCHI conference on Human factors in computing systems*. ACM Press/Addison-Wesley Publishing Co., 43–50.
- [37] Bongshin Lee, Rubaiat Habib Kazi, and Greg Smith. 2013. SketchStory: Telling more engaging stories with data through freeform sketching. *IEEE Transactions on Visualization and Computer Graphics* 19, 12 (2013), 2416–2425.
- [38] Sang Won Lee, Yujin Zhang, Isabelle Wong, Yiwei Yang, Stephanie D. O’Keefe, and Walter S. Lasecki. 2017. SketchExpress: Remixing Animations for More Effective Crowd-Powered Prototyping of Interactive Interfaces. In *UIST ’17*. 817–828. <https://doi.org/10.1145/3126594.3126595>
- [39] Fabrice Matulic, Lars Engeln, Christoph Träger, and Raimund Dachselt. 2016. Embodied interactions for novel immersive presentational experiences. In *Proceedings of the 2016 CHI Conference Extended Abstracts on Human Factors in Computing Systems*. ACM, 1713–1720.
- [40] NewscastStudio. 2017. Weather Solutions: 4 of today’s leading weather software systems. <https://www.newscaststudio.com/2017/06/21/tv-weather-solutions/>.
- [41] Ken Perlin. 2016. The Coming Age of Computer Graphics. In *Sanders Series Lecture*. TUX.
- [42] Ken Perlin, Zhenyi He, and Karl Rosenberg. 2018. Chalktalk : A Visualization and Communication Language – As a Tool in the Domain of Computer Science Education. *ArXiv e-prints* (Sept. 2018).
- arXiv:cs.HC/1809.07166
- [43] Mitchel Resnick, John Maloney, Andrés Monroy-Hernández, Natalie Rusk, Evelyn Eastmond, Karen Brennan, Amon Millner, Eric Rosenbaum, Jay Silver, Brian Silverman, et al. 2009. Scratch: programming for all. *Commun. ACM* 52, 11 (2009), 60–67.
- [44] Thad Starner, Joshua Weaver, and Alex Pentland. 1998. Real-time american sign language recognition using desk and wearable computer based video. *IEEE Transactions on pattern analysis and machine intelligence* 20, 12 (1998), 1371–1375.
- [45] Christian Steins, Sean Gustafson, Christian Holz, and Patrick Baudisch. 2013. Imaginary devices: gesture-based interaction mimicking traditional input devices. In *Proceedings of the 15th international conference on Human-computer interaction with mobile devices and services*. ACM, 123–126.
- [46] Daniel Wigdor, Hrvoje Benko, John Pella, Jarrod Lombardo, and Sarah Williams. 2011. Rock & rails: extending multi-touch interactions with shape gestures to enable precise spatial manipulations. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. ACM, 1581–1590.
- [47] Nora S. Willett, Wilmot Li, Jovan Popovic, and Adam Finkelstein. 2017. Triggering Artwork Swaps for Live Animation. In *UIST ’17*. 85–95. <https://doi.org/10.1145/3126594.3126596>
- [48] Andrew D Wilson and Hrvoje Benko. 2010. Combining multiple depth cameras and projectors for interactions on, above and between surfaces. In *Proceedings of the 23rd annual ACM symposium on User interface software and technology*. ACM, 273–282.
- [49] Andrew D Wilson, Shahram Izadi, Otmar Hilliges, Armando Garcia-Mendoza, and David Kirk. 2008. Bringing physics to the surface. In *Proceedings of the 21st annual ACM symposium on User interface software and technology*. ACM, 67–76.
- [50] Yukang Yan, Chun Yu, Xiaojuan Ma, Xin Yi, Ke Sun, and Yuanchun Shi. 2018. VirtualGrasp: Leveraging Experience of Interacting with Physical Objects to Facilitate Digital Object Retrieval. In *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*. ACM, 78.
- [51] Qi Yang and Georg Essl. 2015. Representation-Plurality in Multi-Touch Mobile Visual Programming for Music. In *NIME 2015*. The School of Music and the Center for Computation and Technology (CCT), Louisiana State University, Baton Rouge, Louisiana, USA, 369–373. <http://dl.acm.org/citation.cfm?id=2993778.2993872>
- [52] Bo Zhu, Michiaki Iwata, Ryo Haraguchi, Takashi Ashihara, Nobuyuki Umetani, Takeo Igarashi, and Kazuo Nakazawa. 2011. Sketch-based dynamic illustration of fluid systems. In *ACM Transactions on Graphics (TOG)*, Vol. 30. ACM, 134.